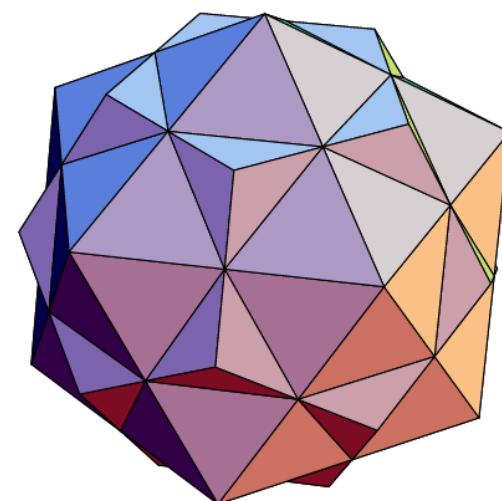


Coordinating web3 computing workloads

Mike Voronov
twitter.com/@vms11



Agenda

- Algorithms in distributed networks
- AquaVM: pi-calculus based approach
- Fluence Labs implementation
- Distributed Choreography and Composition Working Group (**DCC-WG**)

Algorithms in distributed networks

Assumption & Question

Let's say we have a way to run a computation on a peer in the form of function and (possibly stateful and even effectful) services on peers:

$$f_{peer} : \mathbf{CID} \rightarrow \mathbf{CID}$$

How to compose a set of such functions possibly located on different peers?

$$f_{system} = (f_{peer_1} \circ \dots \circ f_{peer_n}) : \mathbf{CID} \rightarrow \mathbf{CID}$$

Qs for f_{system}

- How to put this function where it needs to be?
- How to find a peer ready to execute my code? Move data to code, code to data, etc?
- How to organize many peers to take their parts in a workflow?
- How to organize the network of different peers with different code there?
- How to protect from common attacks, like replay, Sybil and Eclipse?
- How to do all of this in the most efficient way, improve latency, etc?

Approaches



Centralized computing

- Middlemen (cloud, admin)
- No redundancy
- Unverifiable
- No consensus
- Cheap

Decentralized computing

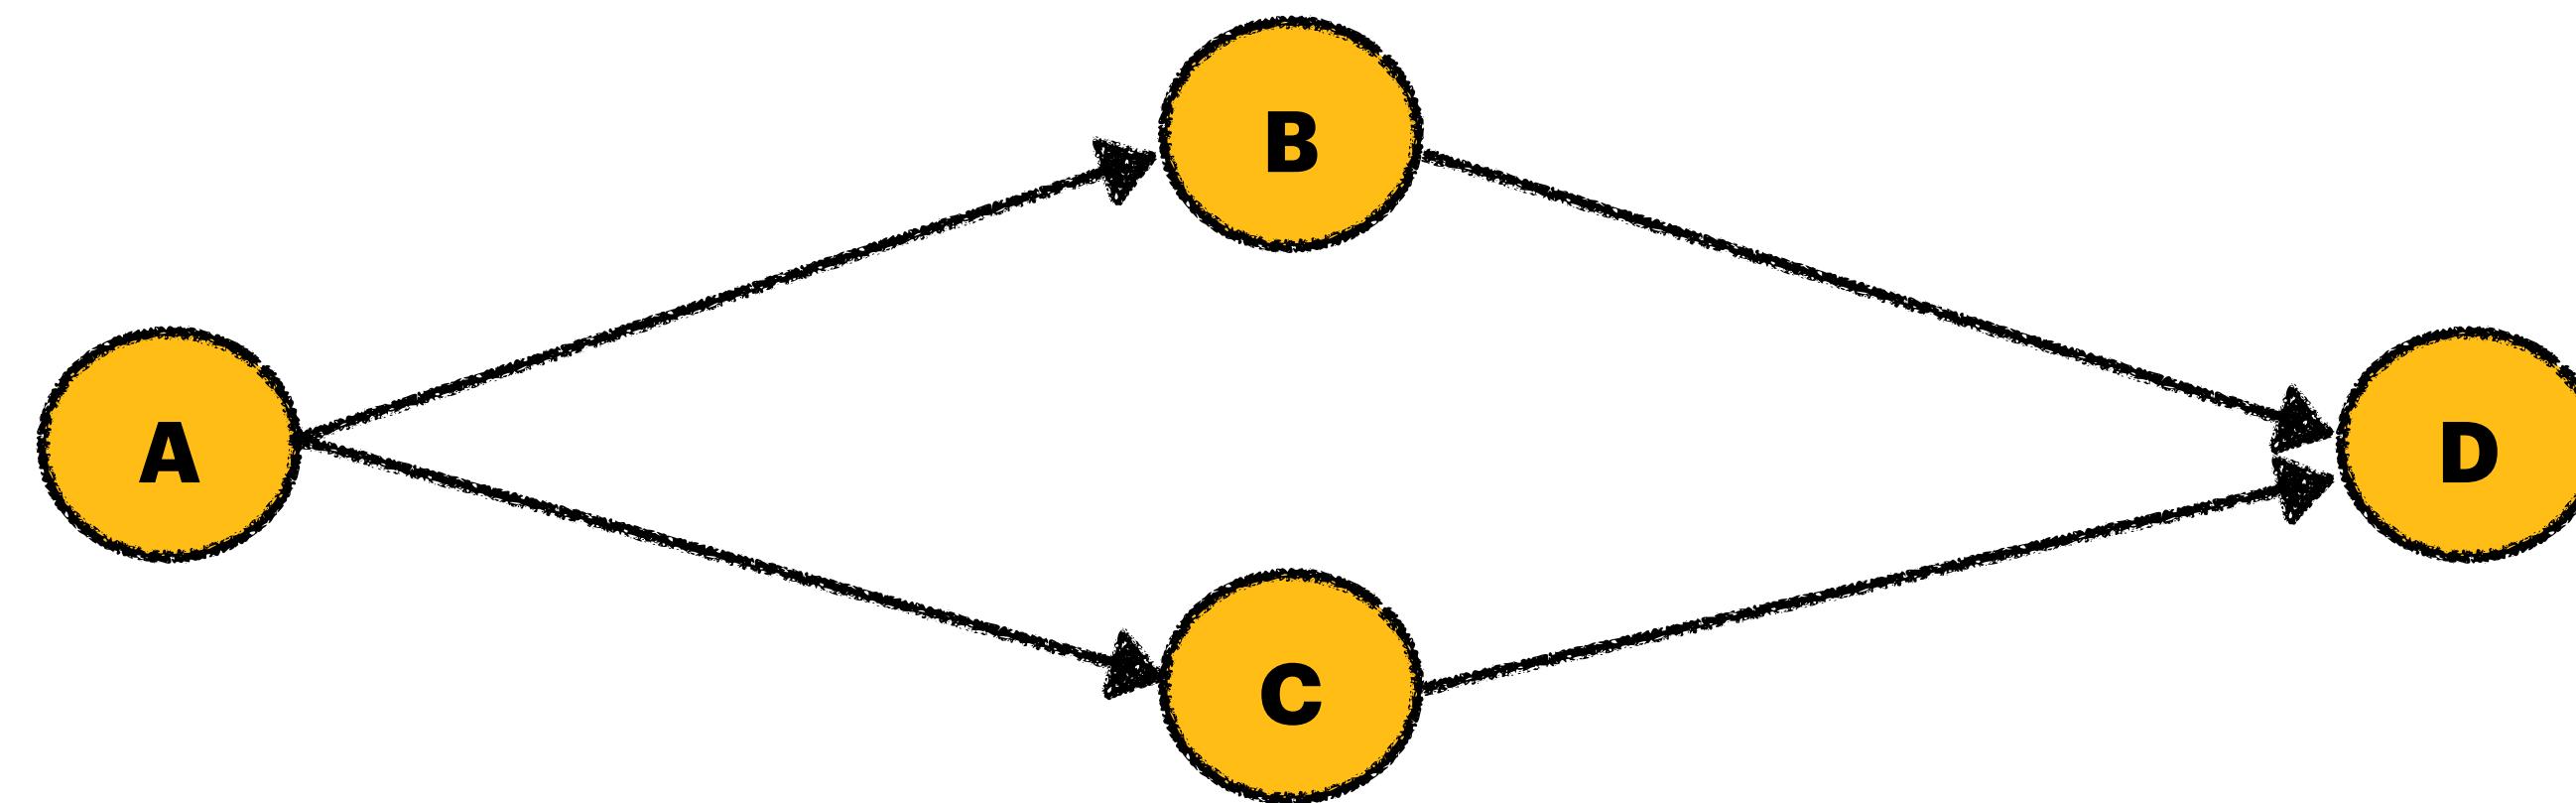
- ✓ No middlemen
- ✓ No redundancy
- ✓ Verifiable (flexible)
- ✓ Optional consensus
- ✓ Cheap

On-chain computing

- No middlemen
- Redundancy
- Verifiable
- BFT consensus
- Expensive

Observation

- there are some models to express and orchestrate decentralized computing, but many of them are expressed with the **fork-join pattern** (such as OpenMP, map-reduce and many others)



- is there any suitable theory for that?

pi-calculus to express fork-join

- Output, Input: $a!b$, $a?b$
- Combinators: $P.Q$, $P+Q$, $P|Q$
- Match, Mismatch: $(x == y)P$, $(x != y)P$
- Restriction: $(\text{new } x)P$
- Replication: $!P := P \mid !P$
- Null: $.0$



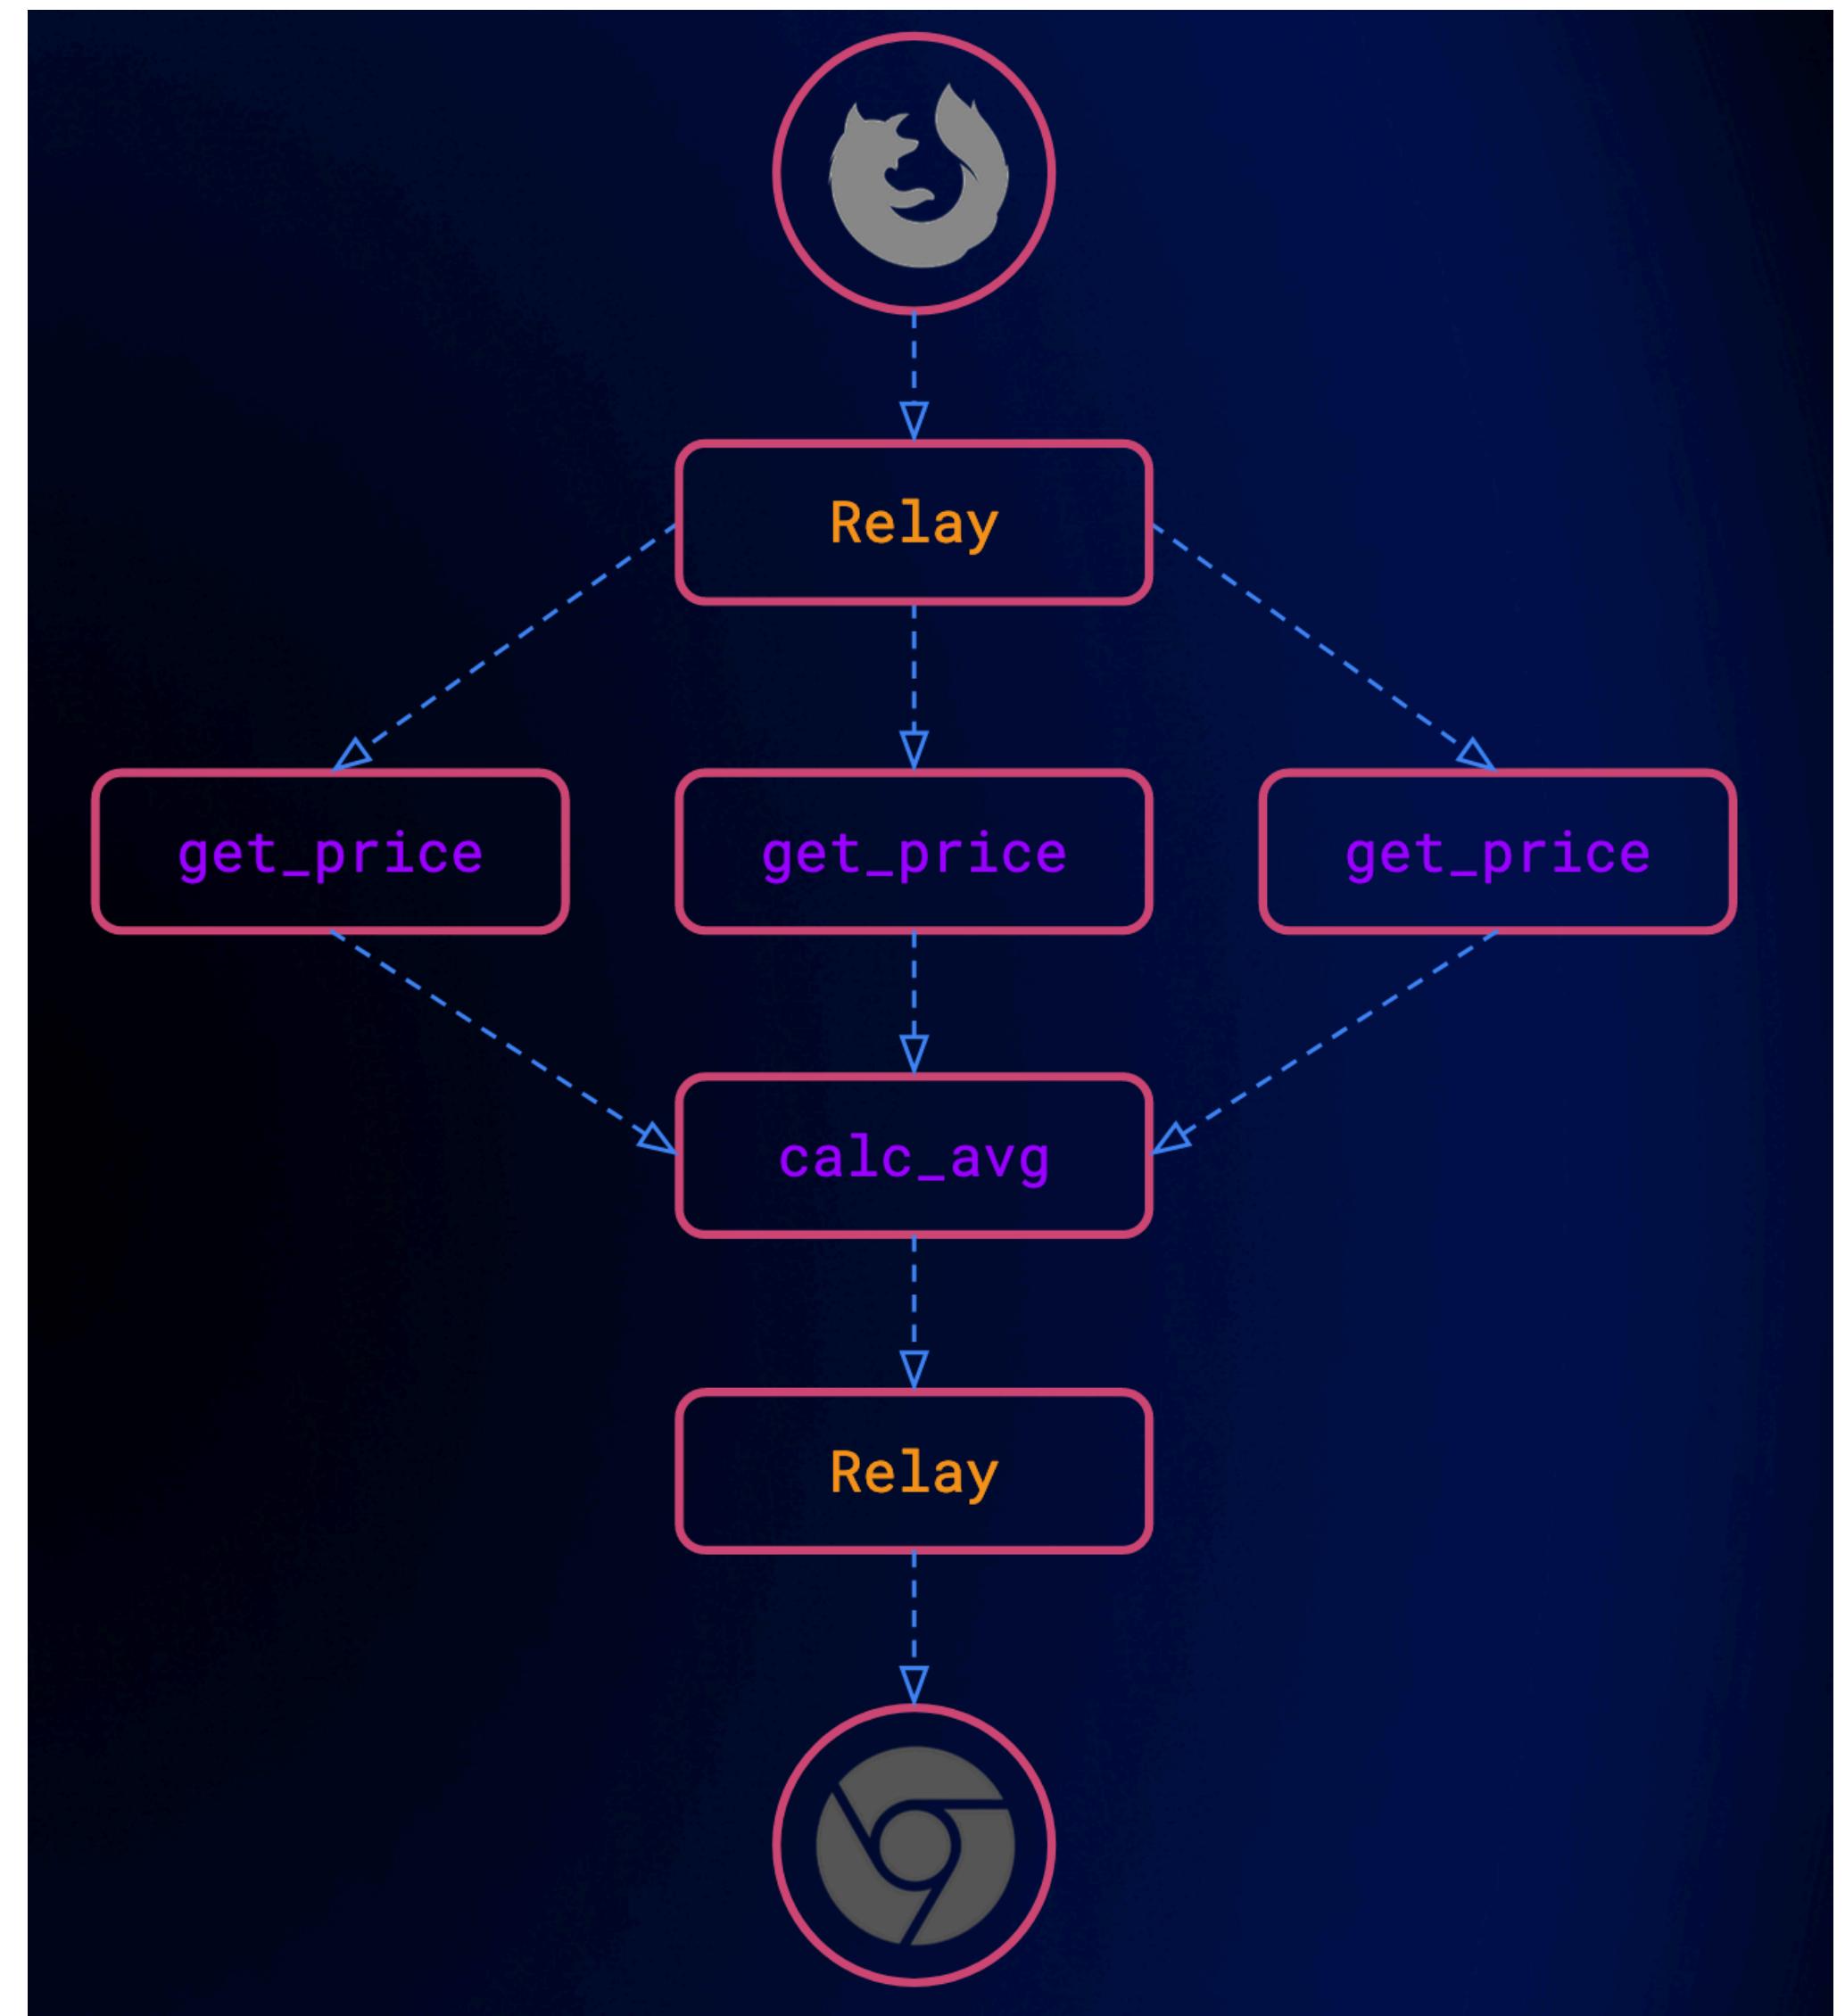
Aqua fork-join example

```
map_reduce.aqua

for p <- peers par:
    prices <- PriceService.get_price()

on avg_peer:
    avg <- AvgService.calc_avg(prices)

on received_peer via relay:
    ReceiveService.receive_avg(avg)
```



AquaVM: pi-calculus based approach

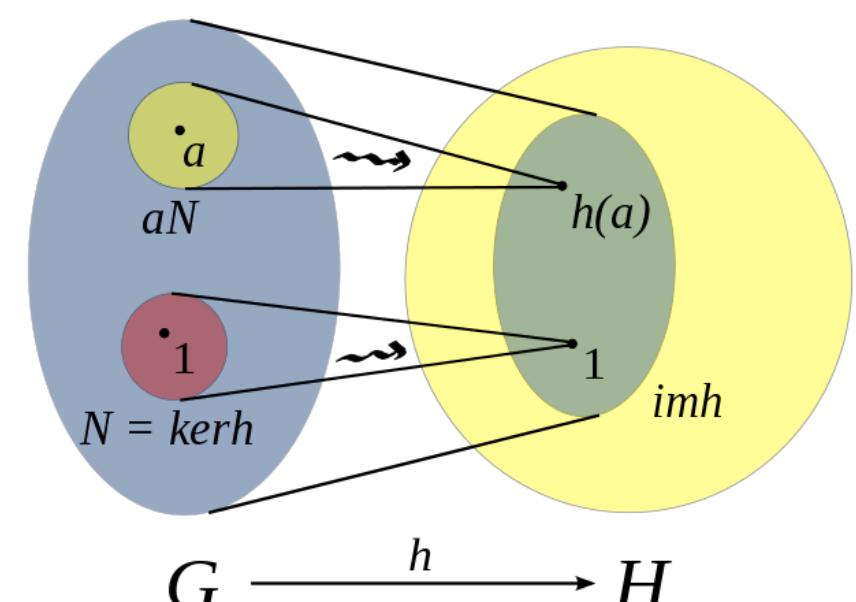
AquaVM's way to rule the complexity

To rule the complexity of development and verify our research ideas we use the idea of homomorphic mapping

$$f: X \rightarrow Y$$

where X is AquaVM

$$f(x + y) = f(x) * f(y)$$



What is Y?

pi calculus

lambda calculus

abstract algebra

category theory

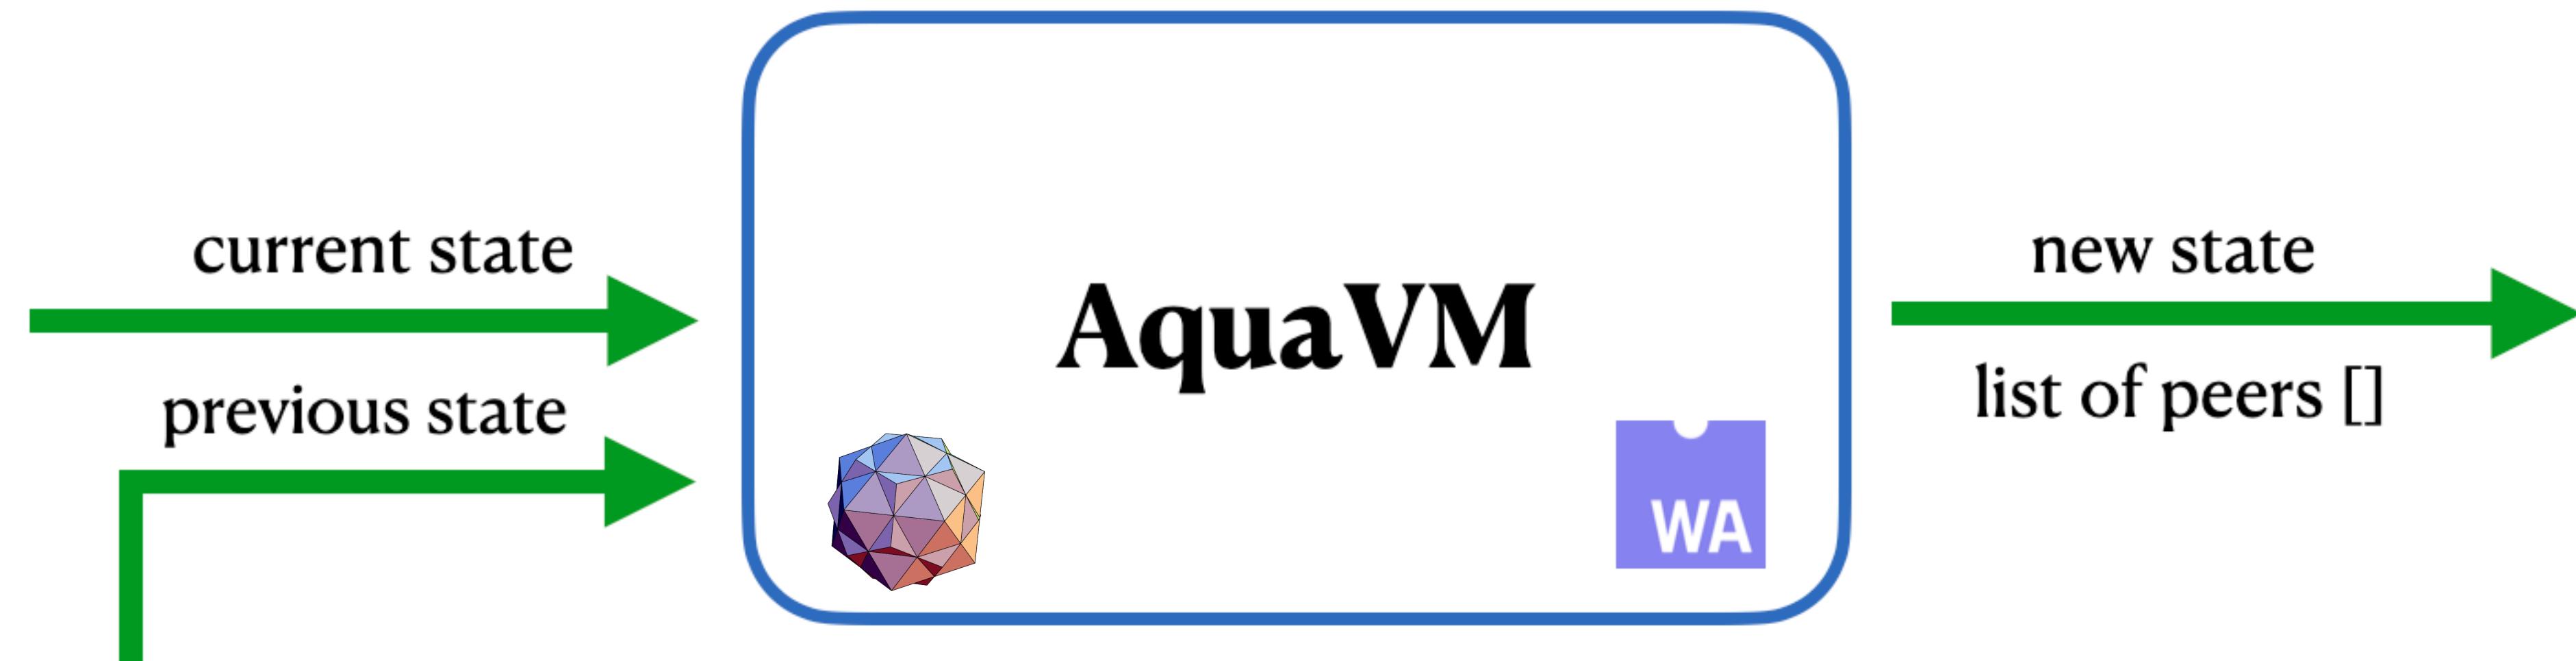
graph theory

automata theory

compiler theory

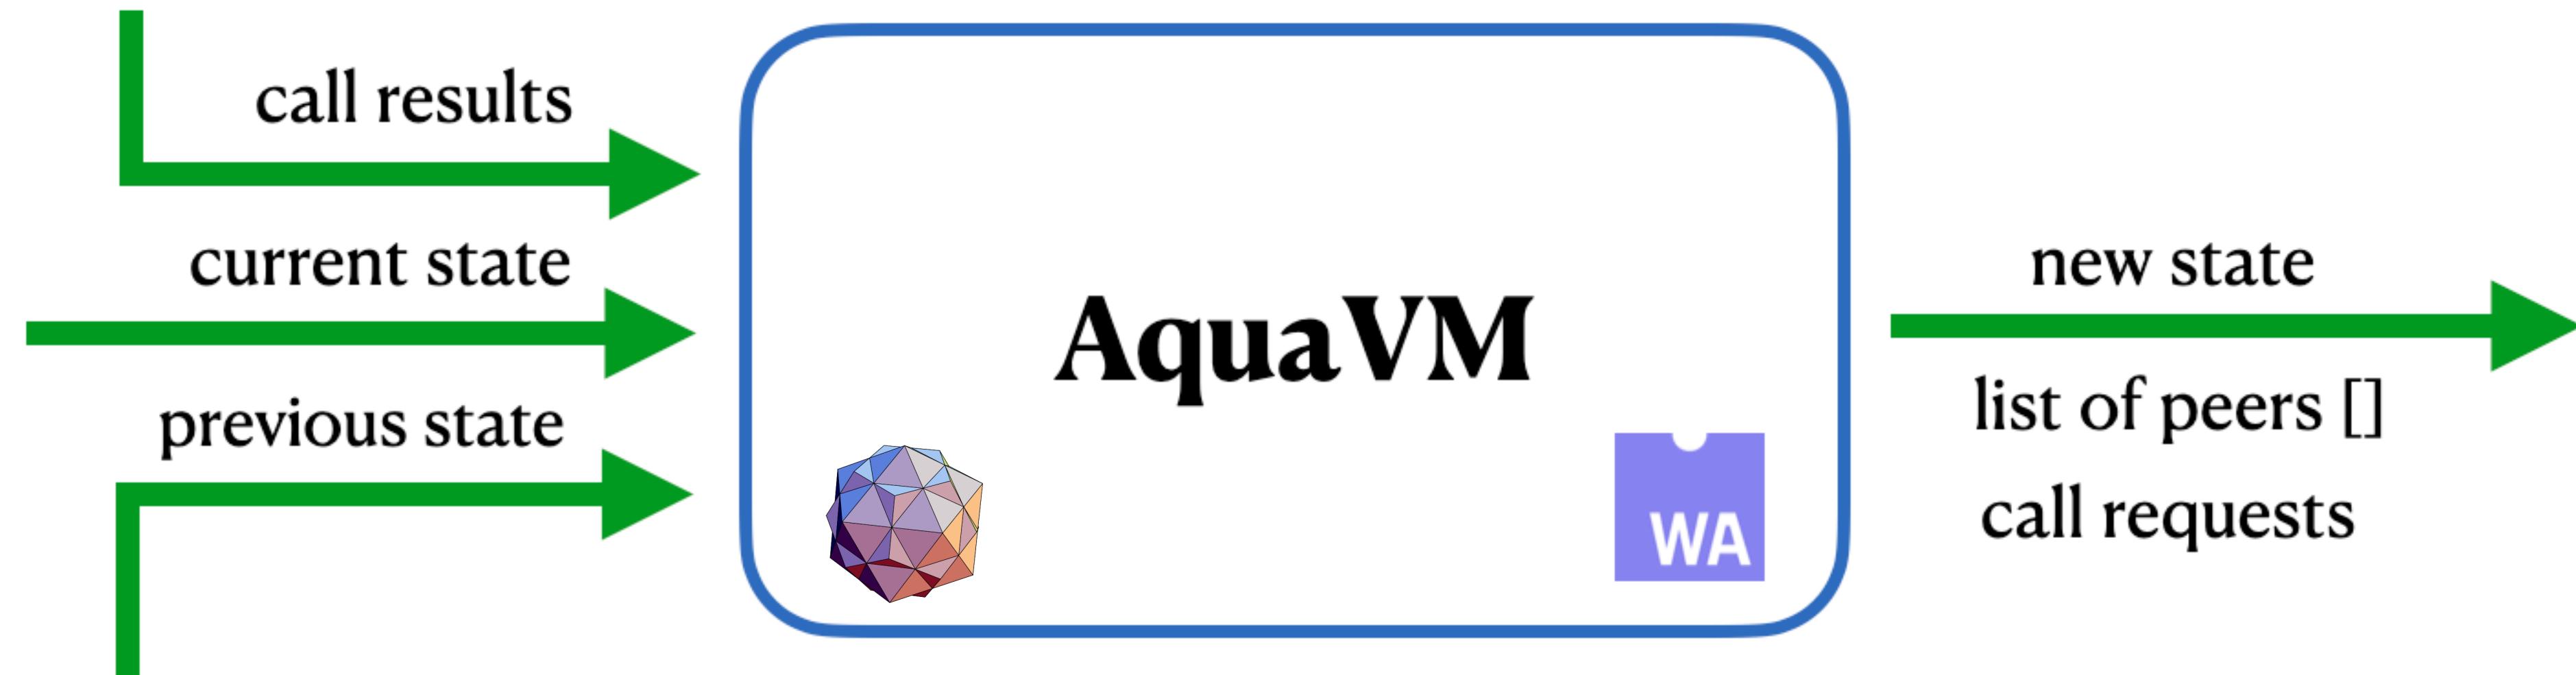
TLA+

AquaVM as FSM



- AquaVM is a pure state transition function
- It takes previous state and state from a particle
 - returns a new state and a list of peers to send this state

AquaVM as FSM



- AquaVM collects parameters of **service calls** that could be executed on this peer and return them after finishing execution
- then it expects to obtain **results** of their execution back

AquaVM: triggers

AquaVM triggers:

- handle data from a new network packet (with the same particle_id)
- handle a call result from a called service

Pros of this scheme:

- interpreter is a pure function
- allows async/parallel service execution on a peer
- execution takes a fixed time

AquaVM instructions

call - call f_peer

AquaVM instructions

call - call f_peer

seq - sequential execution
par - parallel execution
xor - execution with catching errors

AquaVM instructions

call - call f_peer

match/mismatch branching

new - create a new scope
ap - applying a functor
canon - canonicalization

seq - sequential execution
par - parallel execution
xor - execution with catching errors

null - does nothing
never - stops execution
fail - throws an error

fold/next - iteration

category theory

compiler theory

pi calculus

AquaVM: execution example

- every time execution starts from the very beginning of each script
- result state contains a linearized execution trace
- let's consider how this simple fork-join example will be executed by AquaVM:

```
(seq
  (par
    (call "provider_1" ("prices" "get") [] price_1)
    (call "provider_2" ("prices" "get") [] price_2)
  )
  (seq
    (call "average" ("average_service" "calc") [price_1 price_2] avg_price)
    (call "user" ("recieve_service" "show") [avg_price])
  )
)
```

AquaVM: execution example

Let's say that we start execution on a peer with peer_id "user"

```
→ (seq
    (par
        (call "provider_1" ("prices" "get") [] price_1)
        (call "provider_2" ("prices" "get") [] price_2)
    )
    (seq
        (call "average" ("average_service" "calc") [price_1 price_2] avg_price)
        (call "user" ("recieve_service" "show") [avg_price])
    )
)
```

peer_id: "user"

seq takes two instructions and executes the second iff the first was executed

AquaVM: execution example

```
(seq  
→  (par peer_id: "user"  
      (call "provider_1" ("prices" "get") [] price_1)  
      (call "provider_2" ("prices" "get") [] price_2)  
    )  
    (seq  
      (call "average" ("average_service" "calc") [price_1 price_2] avg_price)  
      (call "user" ("recieve_service" "show") [avg_price])  
    )  
)
```

par takes two instructions and executes them not depending on each other

AquaVM: execution example

AquaVM finishes with **next_peers = ["provider_1", "provider_2"]** and **new state**

```
(seq
  (par
    (call "provider_1" ("prices" "get") [] price_1)
    (call "provider_2" ("prices" "get") [] price_2)
  )
  (seq
    (call "average" ("average_service" "calc") [price_1 price_2] avg_price)
    (call "user" ("recieve_service" "show") [avg_price])
  )
)
```

→ peer_id: "user"

call executes function from a service on particular peer with the provided set of arguments bounding result to specified variable name

AquaVM: execution example

AquaVM finishes with **next_peers = ["average"]** and **new state**

```
(seq
  (par
    (call "provider_1" ("prices" "get") [] price_1)
    (call "provider_2" ("prices" "get") [] price_2)
  )
  (seq
    (call "average" ("average_service" "calc") [price_1 price_2] avg_price)
    (call "user" ("recieve_service" "show") [avg_price])
  )
)
```

→ **peer_id: "provider_1", "provider_2"**

call has join behaviour and waits for all their arguments to be defined

AquaVM: execution example

AquaVM finishes with **next_peers = ["user"]** and **new state**

```
(seq
  (par
    (call "provider_1" ("prices" "get") [] price_1)
    (call "provider_2" ("prices" "get") [] price_2)
  )
  (seq
    (call "average" ("average_service" "calc") [price_1 price_2] avg_price)
    (call "user" ("recieve_service" "show") [avg_price])
  )
)
```

→ peer_id: "average"

seq takes two instructions and executes the second iff the first was executed

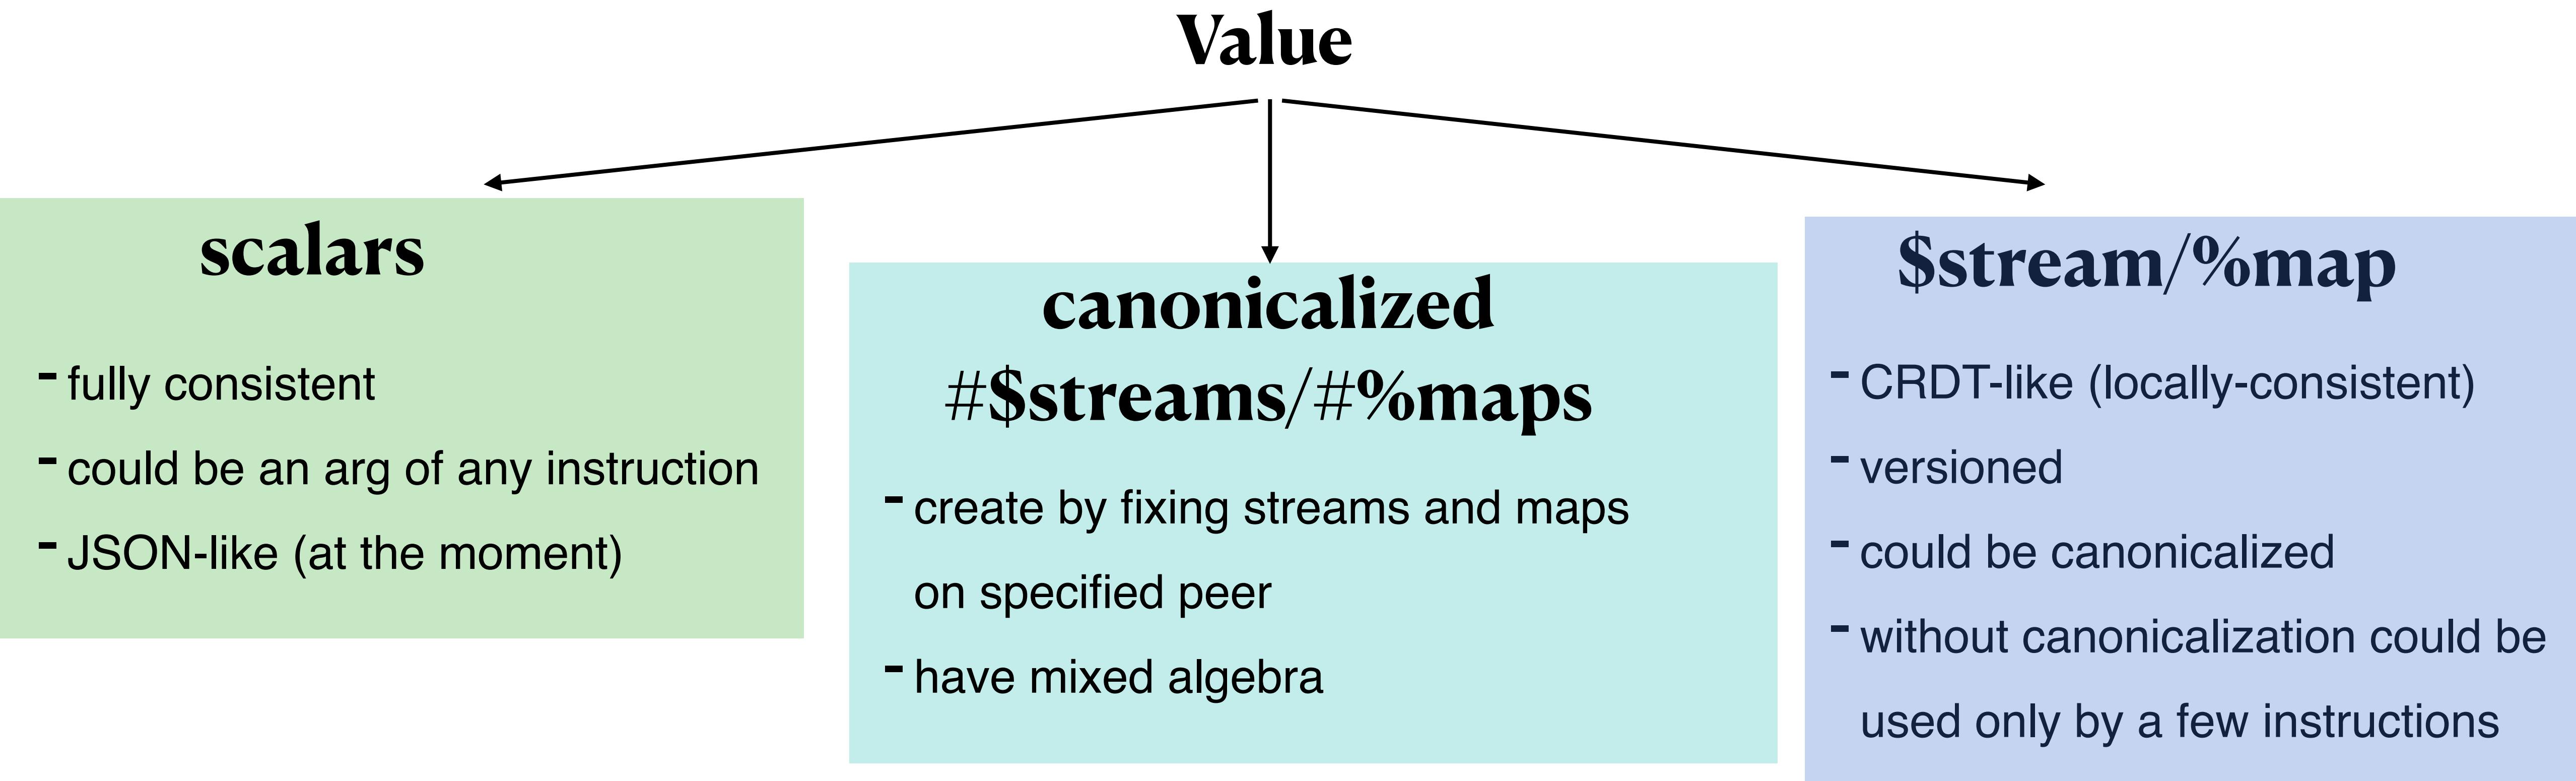
AquaVM: execution example

Finally, `recieve_service.show` is called with the result

```
(seq
  (par
    (call "provider_1" ("prices" "get") [] price_1)
    (call "provider_2" ("prices" "get") [] price_2)
  )
  (seq
    (call "average" ("average_service" "calc") [price_1 price_2] avg_price)
    (call "user" ("recieve_service" "show") [avg_price])
  )
→ )
```

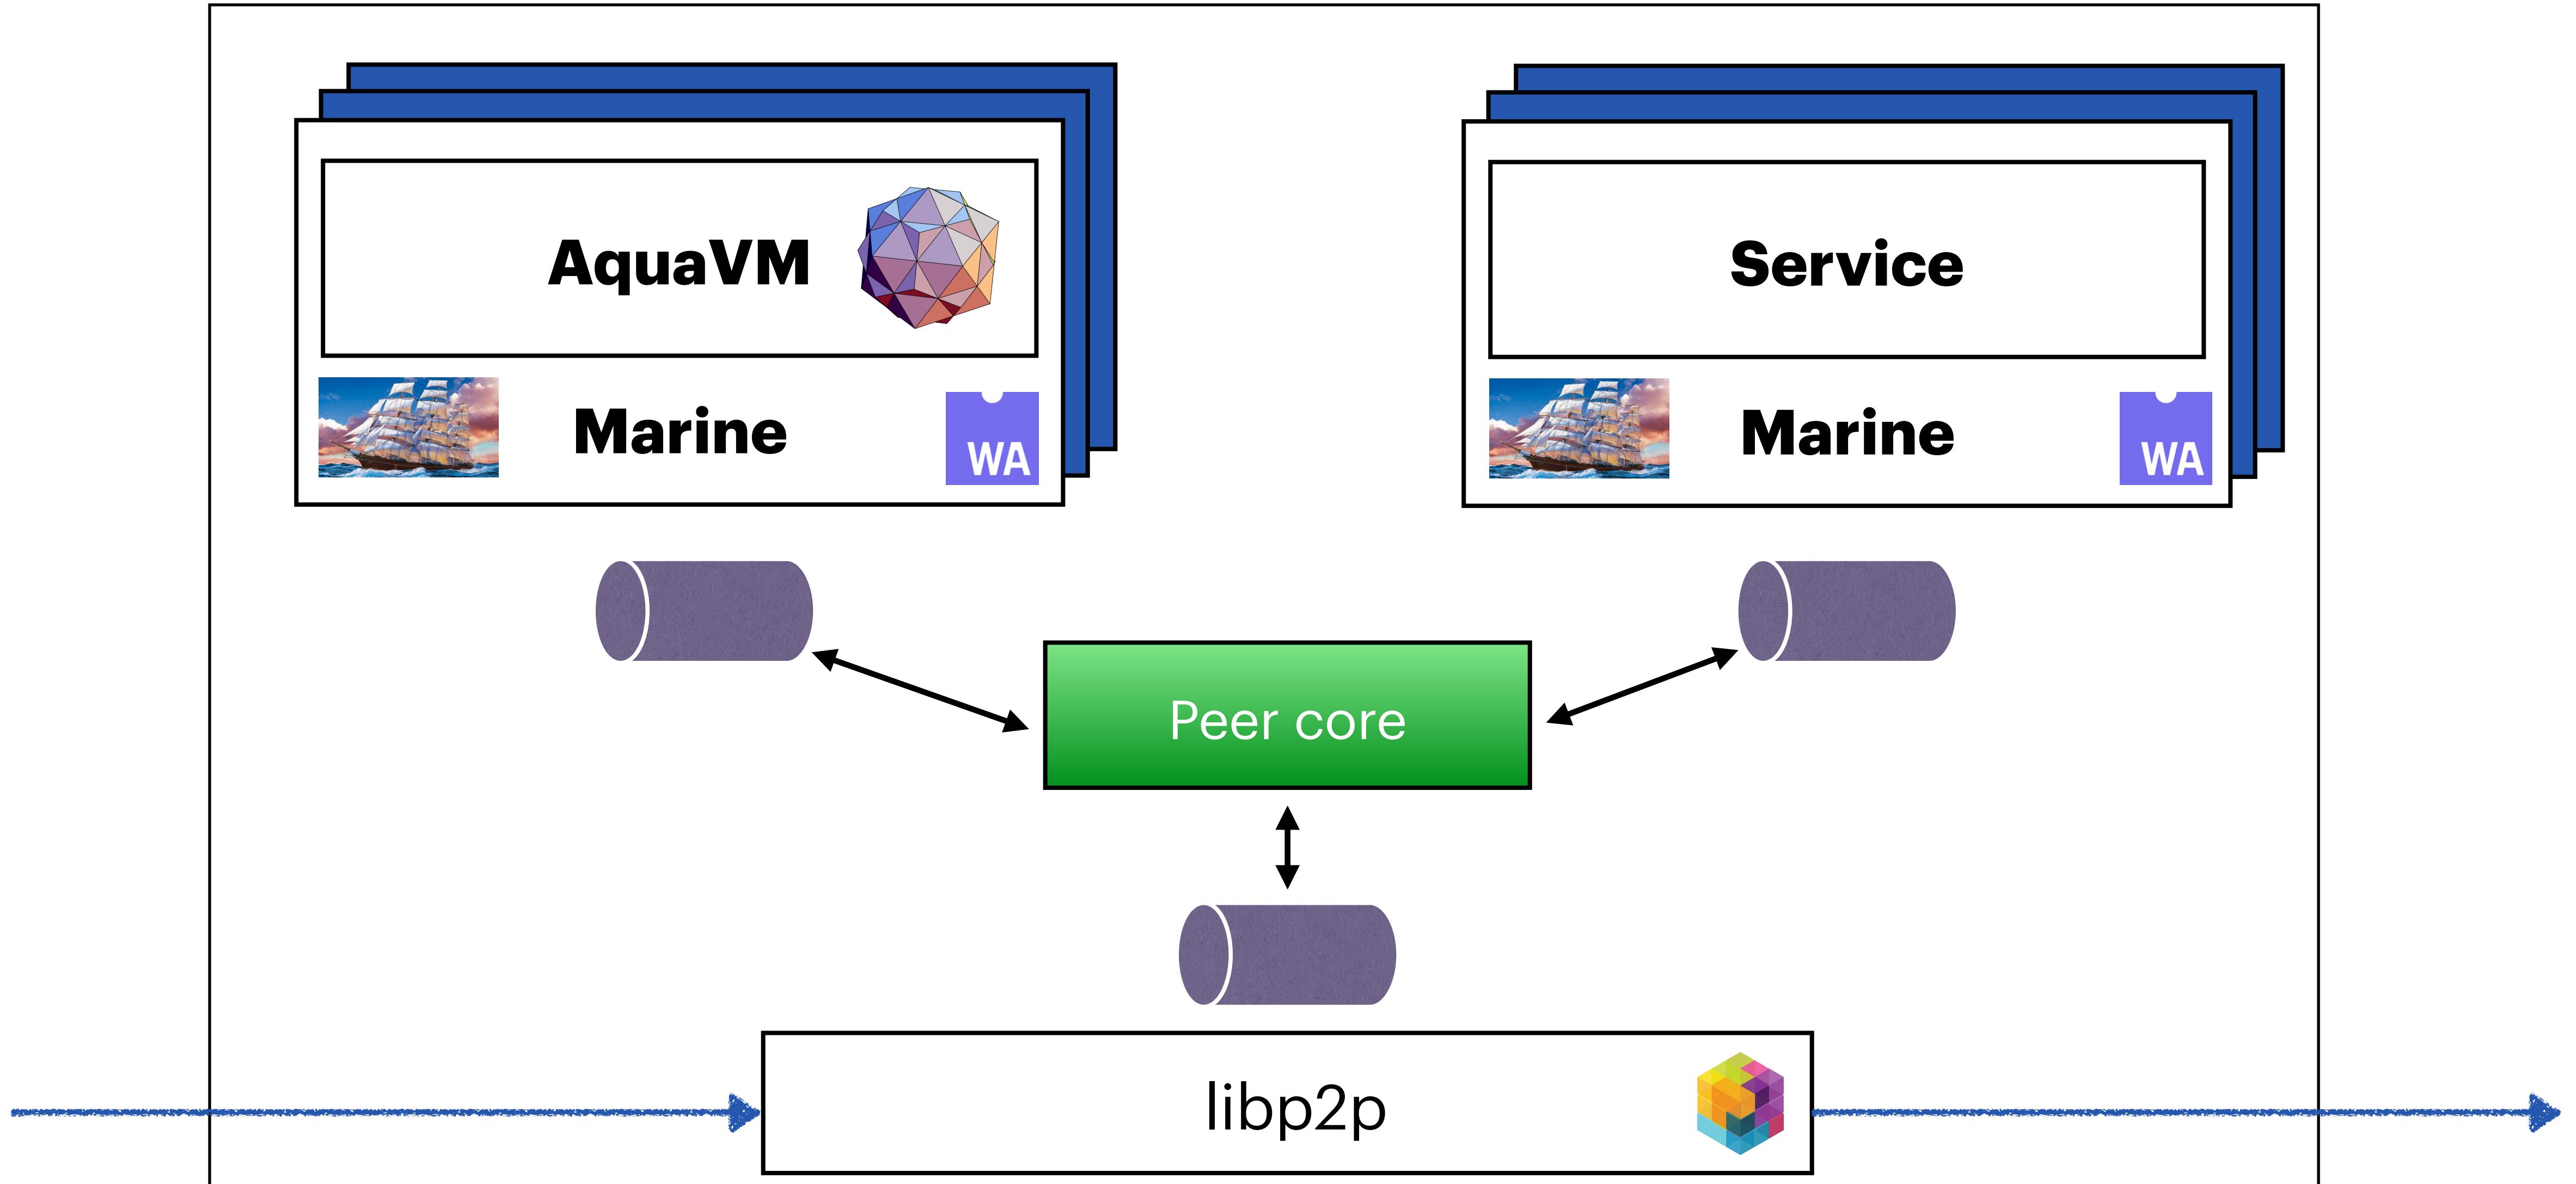
peer_id: "user"

AquaVM: value types



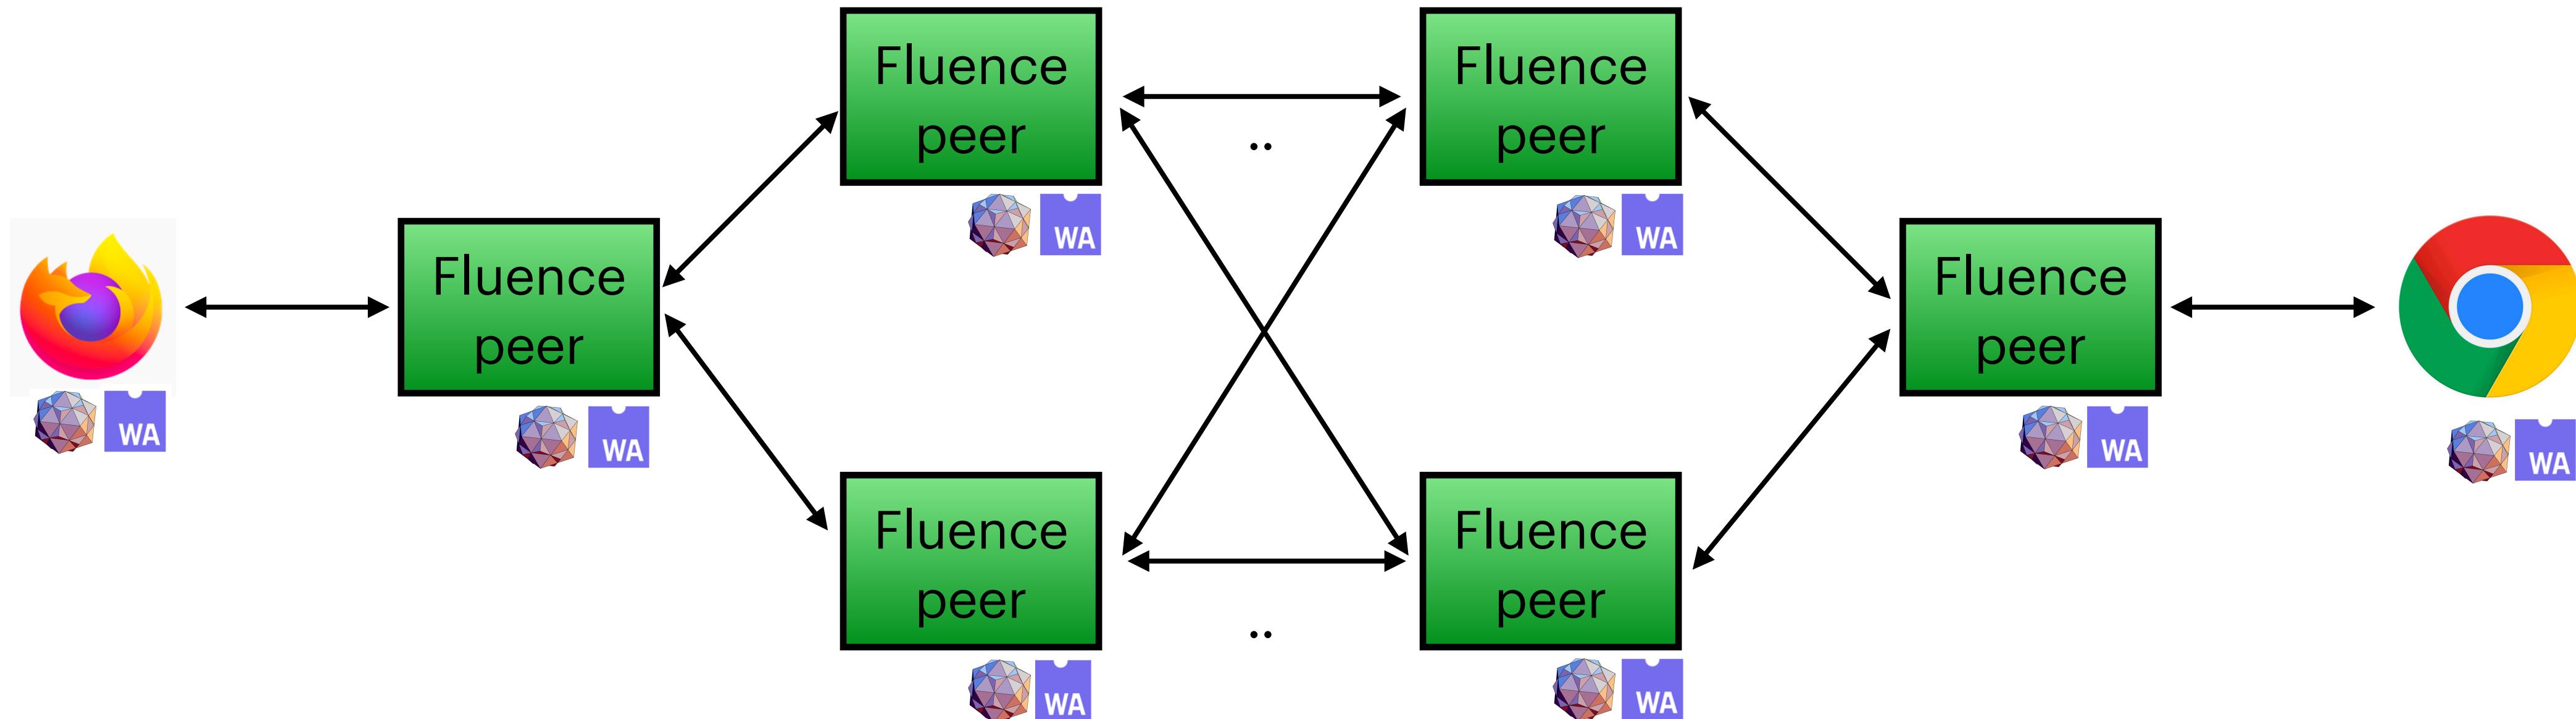
Fluence Labs implementation

Fluence Peer architecture



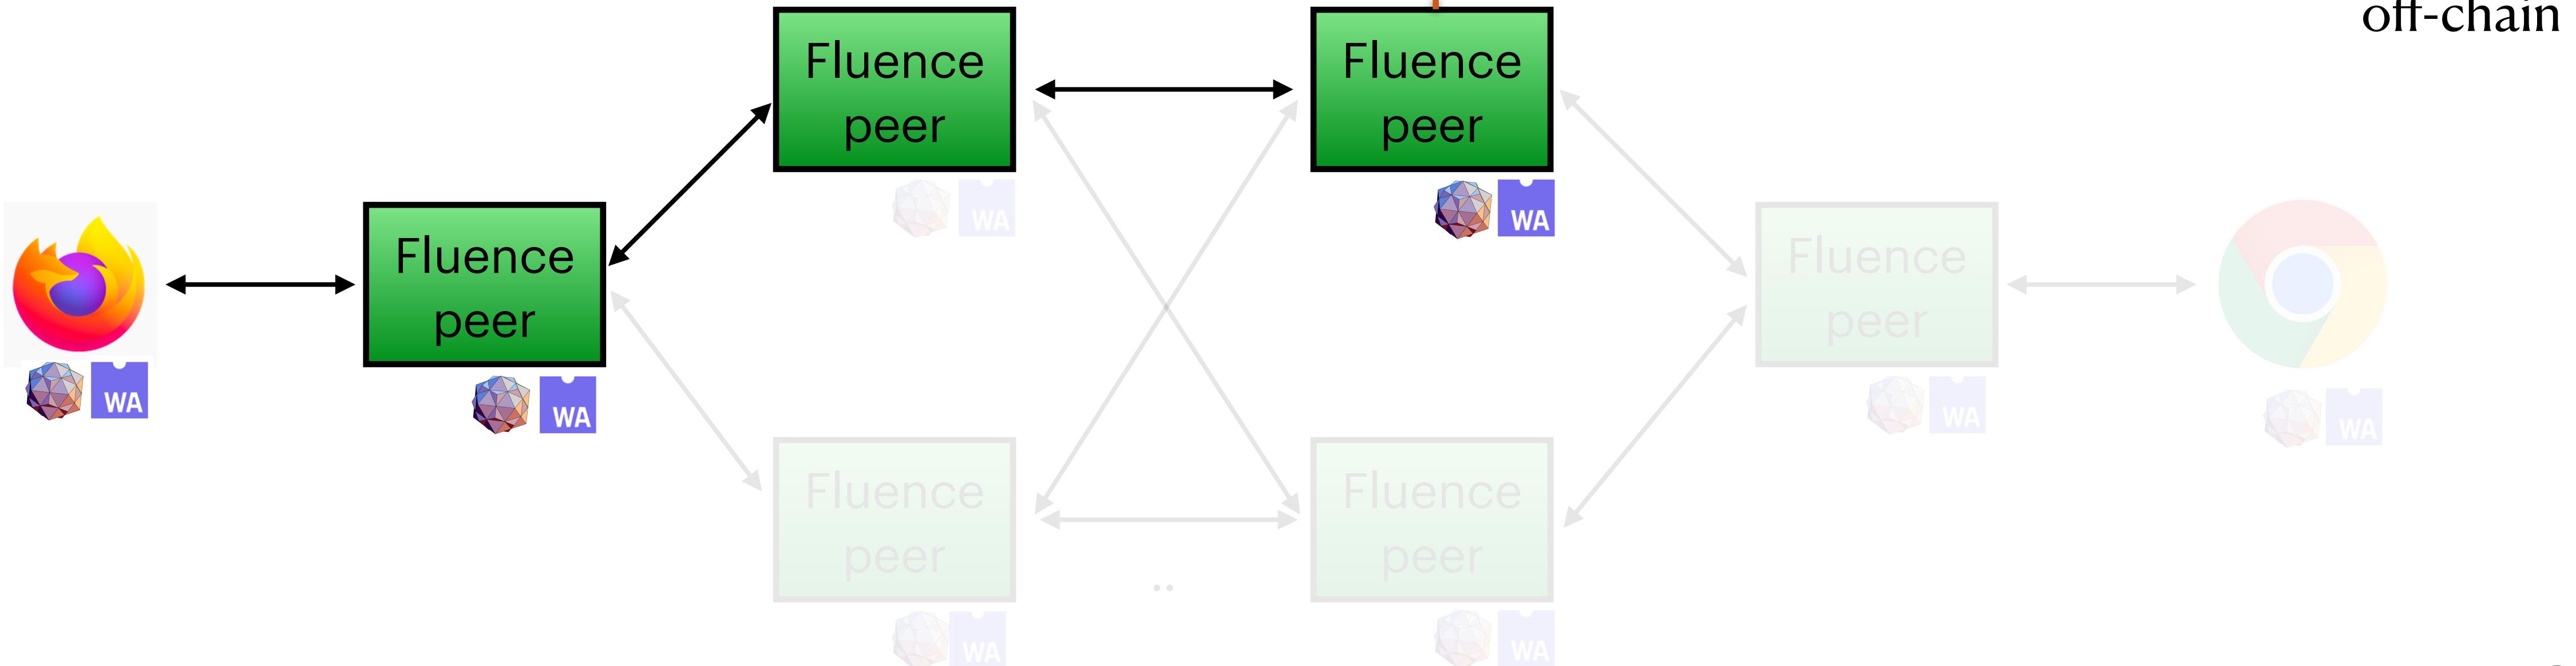
Fluence network

- permissionless
- incentivized
- auditable
- every node is a coordinator
- no implicit consensus



Fluence network

- each peer is incentivized to execute requests and submit a golden network packet to a chain with proofs
- every node participated in forming the packet will be rewarded



DCC working group

DCC WG

Distributed Choreography and Composition Working Group

- **Why**

- DCC is an important component to make distributed and decentralized compute work and enable the use and reuse of (shareable compute) resources
- Provable and independently verifiable DCC implementations are critical to the broad success of decentralized networks

- **What**

- Develop and formalize implementations, best practices and (shareable) processes to enable secure, transparent, provable and verifiable workflow engines, processes and patterns independent of the underlying calculi.
- Tackle challenges such as
 - orchestration models for complex multi-agent permissioned and permissionless p2p networks
 - efficient verification and validation models and processes for both on- and off-chain scenarios
 - formalization approaches over different DCC implementation models for shared reasoning, e.g., via (vanilla) Petri nets or DAGs
 - best practices for payload security and privacy, e.g., e2e

DCC WG: Who Should Care?

Distributed Choreography and Composition Working Group

All Web3 and Web3 stakeholders interested in defining and formalizing a critical part of the decentralized tool and infrastructure stack are welcome and encouraged to participate:

- distributed compute developers
- distributed algorithms developers
- who cares about decentralization of compute

DCC WG



t.me/dcc_wg



github.com/fluencelabs/dcc-wg

my email: mike@fluence.one, tg: @voronovm